

Project Book – Multicore DLX Processor with MESI Coherence Protocol

Project Number: 24-1-1-3024

Student: Yohai Shiloh

ID:

Student: Yarin Koren

ID:

Supervisor: Oren Ganon

Project Carried Out at: Faculty of Engineering, Tel Aviv University

Index

Figures List.....	3
Tables List.....	4
Abstract.....	5
1. Introduction	6
Goals of the Project.....	6
Motivation.....	6
Approach	6
Comparison to Other Multicore Protocols.....	6
2. Theoretical Background	7
Cache	7
The MESI Protocol.....	7
Cache Coherence in Shared Bus Architectures.....	8
Instruction Memory Unit.....	8
General Design Guidelines and Execution Method	8
3. Simulation	9
Environment.....	9
Simulation Objectives.....	9
Key Observations	9
Execution example – Bubble Sort.....	9
Summary of Simulation Coverage:.....	13
4. Implementation	14
4.1 Hardware Description	16
Detailed explanation:.....	17
4.2 Software Description	25
5. Analysis of Results.....	26
Functional Verification	26
Performance Metrics	26
Comparison to other Cache Coherence Protocols: MESI vs MOESI, MSI, and Dragon.....	28
6. Conclusions and Further Work	30
Achievements	30
Improvements Observed	30

Further Work.....	30
Takeaways	30
7. Project Documentation.....	31
User guide	31
8. References.....	32
Appendix A. Inputs outputs tables.....	33
Appendix B. The Simplified DLX Instruction Set Architecture.....	40

Figures List

Figure 1 : System Block Diagram.....	5
Figure 2 : MESI State Diagram.....	7
Figure 3: Example of execution through simulation	9
Figure 4: Multicore architecture's workflow diagram	14
Figure 5 : Multicore architecture Top Level diagram.....	14
Figure 6 : Core architecture diagram	15
Figure 7: Cache structure diagram	15
Figure 8: Cache DataPath architecture diagram	16
Figure 9: Bus Control Unit.....	17
Figure 10: bus control state diagram.....	18
Figure 11: Core Unit.....	18
Figure 12: Fetch State Check Unit.....	18
Figure 13: Cache Unit.....	19
Figure 14: AS_N – DLX control signal DEMUX Unit.....	19
Figure 15: Instruction Memory Unit.....	19
Figure 16: DLX data in MUX Unit.....	20
Figure 17: ACK signal MUX Unit	20
Figure 19: Cache control state diagram.....	21
Figure 18: Cache Control Unit.....	21
Figure 20: Address Calculator Unit	22
Figure 21: Cache DataPath Unit.....	22
Figure 22: Tag Memory flow diagram.....	23
Figure 23: Tag Memory Unit	23
Figure 24: Data In MUX Unit.....	24
Figure 25: Cache Data Memory Unit	24
Figure 26: Programs runtime graph	26
Figure 27: Speedup ratio graph.....	27
Figure 28: Execution example – the array before sorting.....	28
Figure 29: Execution example – the array after sorting	28

Tables List

Table 1: Verification coverage.....	13
Table 2: Run time on each processor for different programs	26
Table 3: Speedup ratio for any two processors.....	27
Table 4: Different coherence protocol comparison.....	29
Table 5: Bus Control Unit Input Output table	33
Table 6: Core Unit Input Output table	33
Table 7: Fetch State Check Unit Input Output table	34
Table 8: Cache Unit Input Output table	34
Table 9: AS_N – DLX control signal DEMUX Unit Input Output table	35
Table 10: Instruction Memory Unit Input Output table.....	35
Table 11: DLX data in MUX Unit Input Output table	35
Table 12: ACK signal MUX Unit Input Output table	36
Table 13: Cache Control Unit Input Output table	36
Table 14: Address Calculator Unit Input Output table	37
Table 15: Cache DataPath Unit Input Output table.....	37
Table 16: Tag Memory Unit Input Output table.....	38
Table 17: Data in MUX Unit Input Output table	39
Table 18: Cache Data Memory Unit.....	39

Abstract

This project extends the classical DLX architecture (which was implemented during the ACSL course) into a dual-core multicore system, addressing performance limitations by implementing parallelism and cache coherence. The project enhances a single-core DLX CPU by adding instruction memory and data caches to each core and introducing a cache coherence mechanism based on the MESI protocol. This ensures data consistency across local caches and enables scalable parallel processing.

A physical implementation was also performed using an FPGA board and the RESA environment, validating both functionality and performance.

The project includes a Verilog based implementation, using both code files and schematic files into a full processor functional design, while supporting the original architecture and instruction set of the DLX processor.

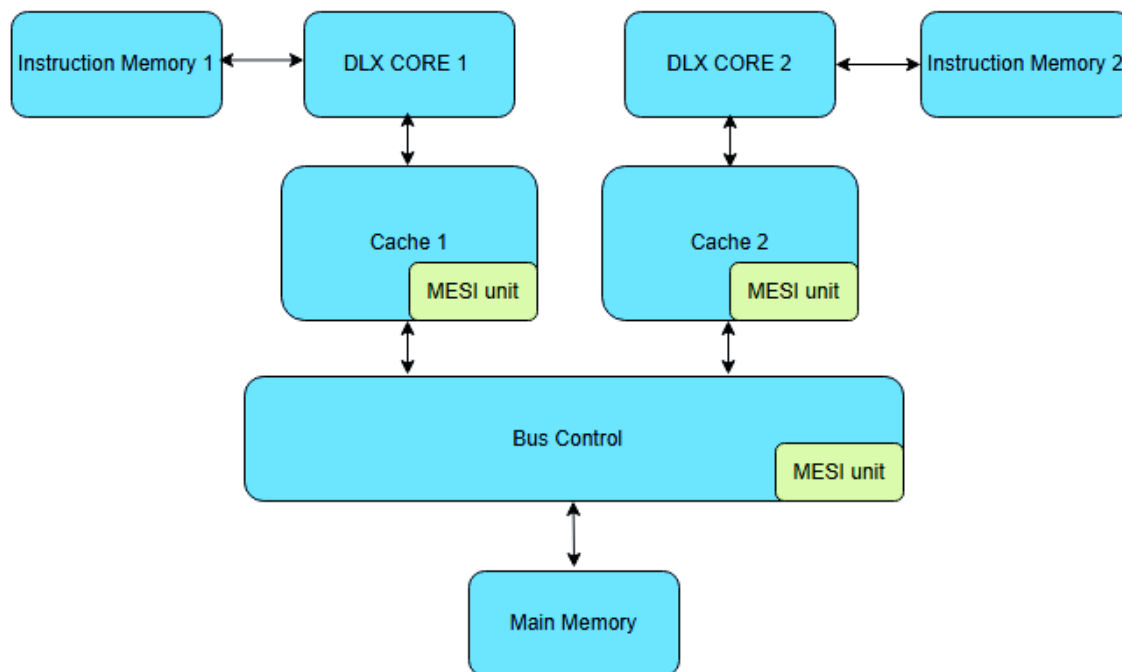


Figure 1 : System Block Diagram

1. Introduction

Goals of the Project

- Extend a single-core DLX processor to a parallel, dual-core system.
- Reduce access time and improve memory bandwidth through per-core instruction and data caches.
- Ensure correctness through the implementation of the MESI coherence protocol.

Motivation

Traditional single-core systems face bottlenecks when attempting to improve performance solely through higher frequencies. Multicore architectures address these challenges by enabling concurrent execution of multiple threads. However, this requires careful management of shared memory and coherence to avoid erroneous behavior, e.g. using an invalid memory that has been modified by other thread, etc.

The use of MESI enables efficient sharing and synchronization of data between caches while reducing unnecessary communication with main memory.

Approach

We reused the basic DLX core structure developed in the ACSL course¹ and expanded it into a multicore platform with the following key additions:

- Independent instruction memory per core to avoid bus contention.
- Data caches for each core.
- A shared bus with arbitration logic.
- A MESI state machine per cache controller.
- A bus monitor to support snooping.

Comparison to Other Multicore Protocols

Whereas some simpler architectures use a shared memory without coherence mechanisms, our design integrates MESI to ensure correctness in data sharing. Compared to directory-based coherence models, MESI provides a balance between complexity and efficiency, suitable for small-scale multicore systems.

More broad comparison is done in section 5.

¹ The Advanced Computer Structure Lab course. During the course, a simplified DLX processor was built by us, under the guide and supervision of the course staff.

Figure 2 : MESI State Diagram

Cache Coherence in Shared Bus Architectures

All caches snoop bus transactions to keep track of changes to shared data. Each cache determines its actions based on the bus signals (BusRd, BusRdX, etc.), enabling decisions about exclusivity.

This passive coherence mechanism avoids the complexity of directory tracking and is ideal for dual-core systems.

Instruction Memory Unit

Each core is assigned a dedicated instruction memory module to avoid contention during instruction fetch. These ROM-based memories are preloaded using custom compilation and initialization tools, enabling both deterministic performance and independent operation of the cores.

General Design Guidelines and Execution Method

The basic engineering point of view was to take the original DLX as a "close box" and build the advanced multicore processor to wrap around it. As a result, the execution flow of the processor includes the following steps:

- The DLX processor fetches an instruction from the instruction memory unit
- If there is an approach to data (e.g. read or write) the DLX addresses the cache to check if the data is available inside the cache.
- If not, the cache will start a transaction to the main memory to import/export data from it
- If the required data is available only in cache of another core (when the block is in the Modified mode), it will be imported from this cache, while simultaneously updating the main memory

3. Simulation

The project included a suite of simulations aimed at verifying functionality and coherence.

Environment

The simulations were written in Verilog and run using Xilinx's Isim platform. Each DLX core ran independently, executing test programs involving shared memory.

The simulation was executed using Verilog based text-bench, performing various

Simulation Objectives

- Validate the functional correctness of the new additions (the cache, instruction memory etc.)
- Validate MESI protocol transitions.
- Ensure correctness of shared data accesses.
- Measure cycle counts and identifies speedup.

Key Observations

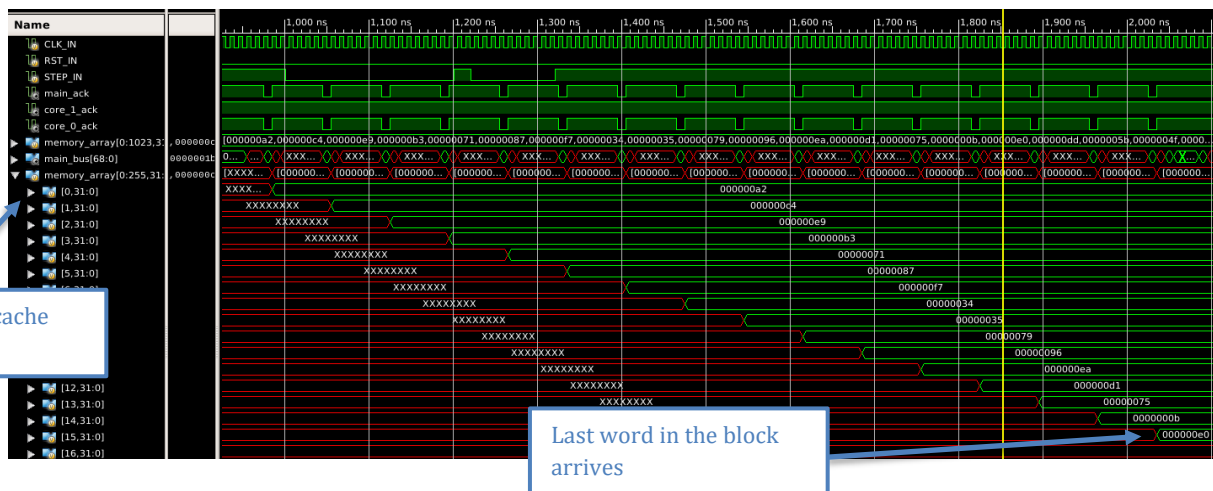
- Cache state transitions behaved as expected.
- The shared line was correctly asserted during shared reads.
- Write invalidation occurred precisely when needed.

Execution example – Bubble Sort

To demonstrate the work of the project, below is attached the execution of one of the programs, the bubble sort program. To avoid excessive length, only key stages are shown, and only main signals and data values are presented.



Loading block into cache No. 0:

Figure 3: Example of execution through simulation

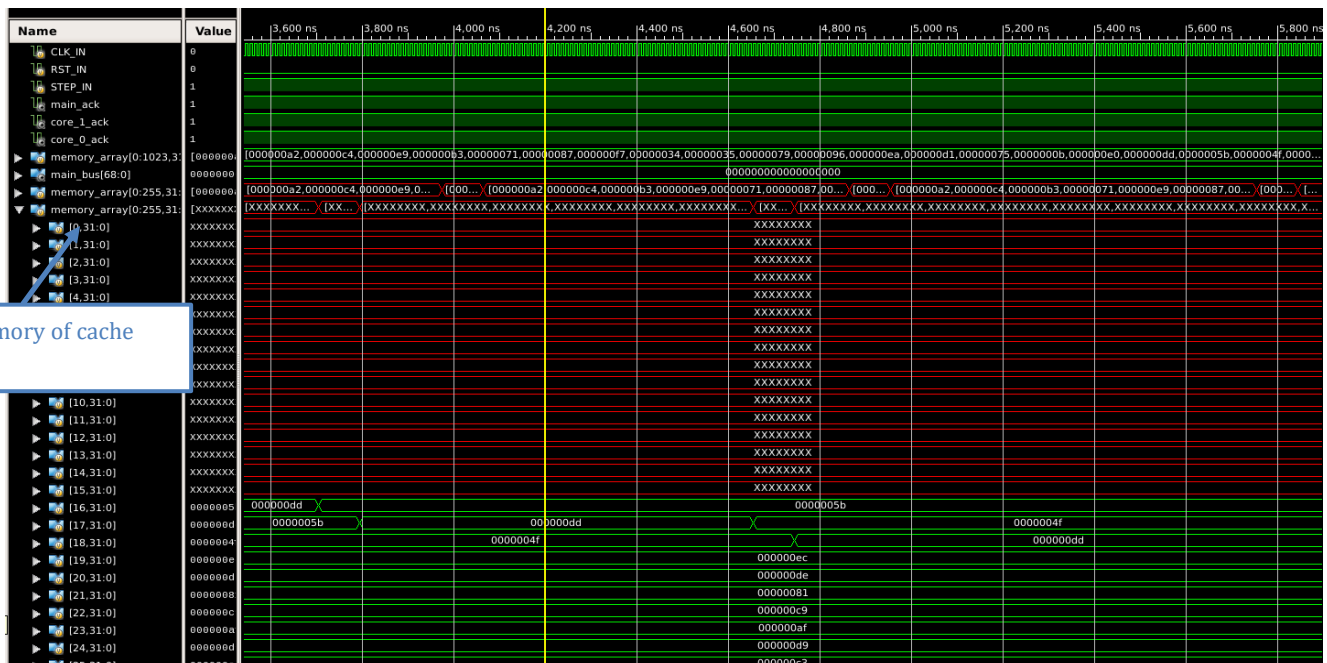


Starting the sort in core 0:


Finding the biggest word(000000f7) and moving it to the last place




Start sorting in core 1(simultaneously):



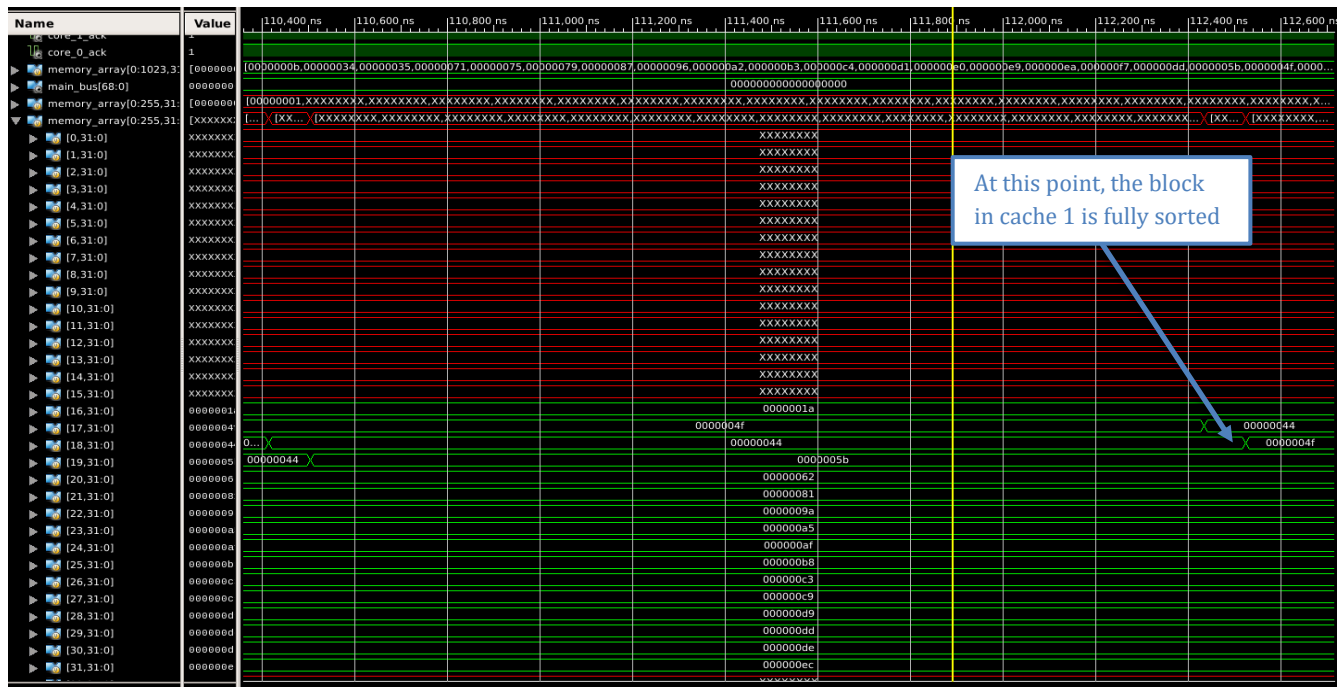
Done the sort in core 0:




Clearing the block from cache No. 0 to the main memory:



Done sorting in core 1:



Merging the two sorted blocks into a fully sorted array and saving it in the main memory (in addresses 48-79):



Summary of Simulation Coverage:

Using the simulation and special Assembly code (see section 4.2), the design was validated using coverage table, a common method to verify the functionality of a system²:

Table 1: Verification coverage

Unit under test	Test/Tests	Tested by -	Notes
Instruction memory unit	Small input handling	Very short program	
	Very large input handling	Very long program	
	Parallel execution by multiple cores from separated units	Parallel programs	
Cache unit	Data transmission from the main memory	Read/Write instructions on missing blocks ("cache miss")	On single-core DLX with cache extension ³
	Data transmission to the core	Read/Write instructions on existing blocks ("cache hit")	On single-core DLX with cache extension
	Updating data in the main memory	Read/Write instructions on missing blocks, while its place in the cache is occupied ("clearing "dirty" block")	On single-core DLX with cache extension
MESI Unit	Transition I to S, E, M	Read/Write instructions	
	Transition S to I	Read/Write was performed by another core	
	Transition S to E, M	Read/Write instructions	
	Transition E, M to I, S	Read/Write was performed by another core	
	Transition E, M to E, M	Read/Write instructions	Both by read/write to the same block and by calling to another block
Bus controller	Correct arbitration between the cores	Various combination of bus usage request from the cores in every possible order	
	Data transmission		
	Signal transmission	IN_INIT, ack_N, wr_N etc. from/to each MESI unit and main memory	
	Snooping	Read/Write instructions	This test is for both bus control and MESI units

² Since the design was made through the engineering point of view to keep the original DLX design as a close box (see chapter 2, section **General Design Guidelines and Execution Method**), its functionality was ensured, and it needed no additional validation as part of the project

³ The simple cache unit was made as part of the first half of the project, which included a single core DLX with cache extension only.

4. Implementation

The design implements the following workflow as the general operation of the processor:

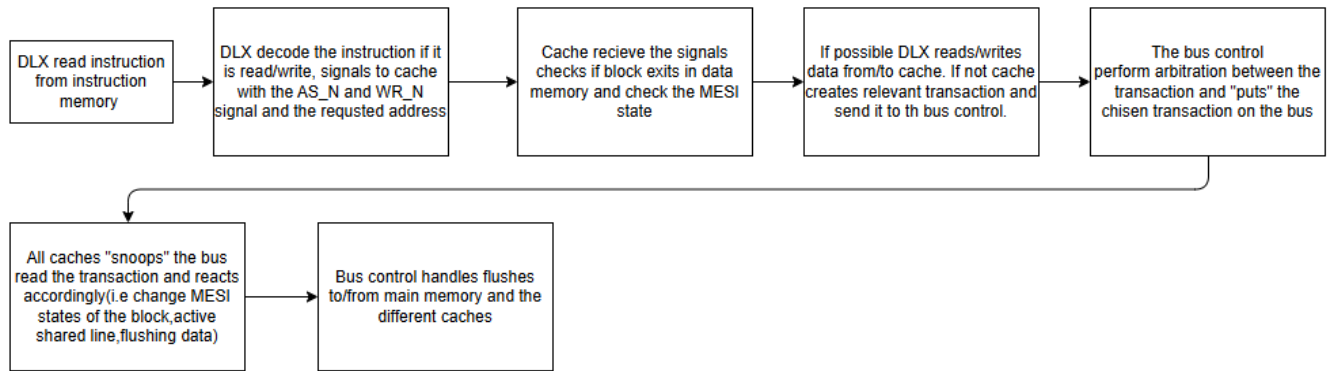


Figure 4: Multicore architecture's workflow diagram

The design structure is made of four main blocks – main memory, bus control unit and two identical processing cores, and are described in the following block diagrams:




Figure 5 : Multicore architecture Top Level diagram

Each core consists of the following structure:

A simple DLX core, a cache memory unit, and supporting logic and small units around the first two, to allow correct data flow and keep the compatibility with the original DLX structure and operation (each component will be discussed in detail in the following sections).

The internal structure of each core is as shown in the following diagram:

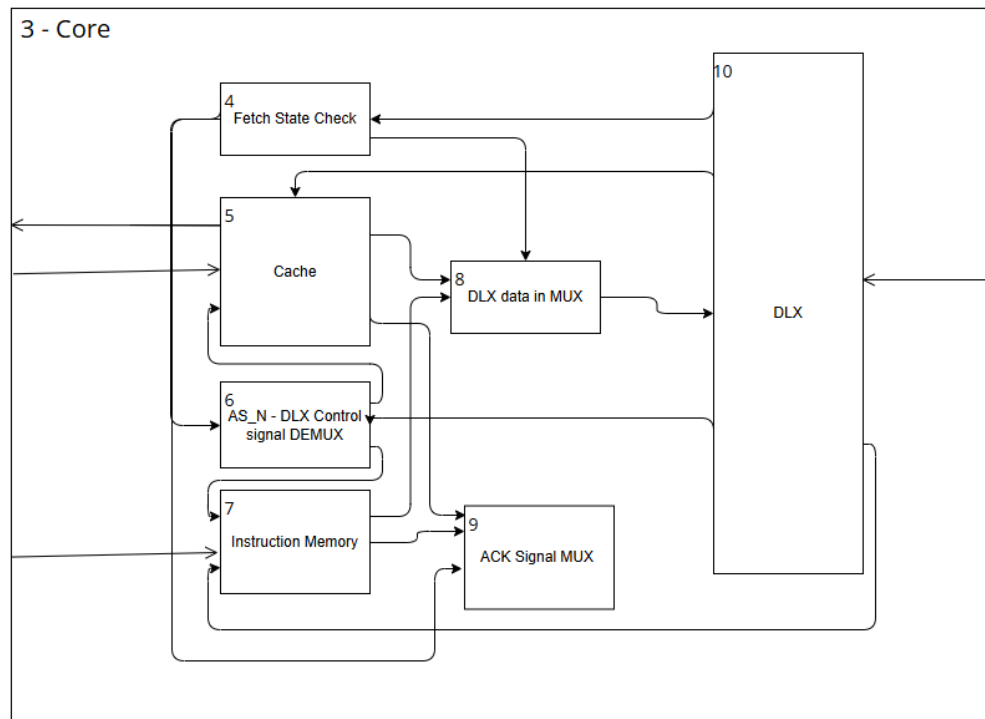


Figure 6 : Core architecture diagram

Each cache memory unit is built of several parts as well – a controller to control its different functions, Cache datapath that store, edit and route the data into and from the cache and address calculator to modify the raw address values the cache receives from the DLX into the cache address form.

The structure of the cache and of the datapath unit is shown in the next diagrams:

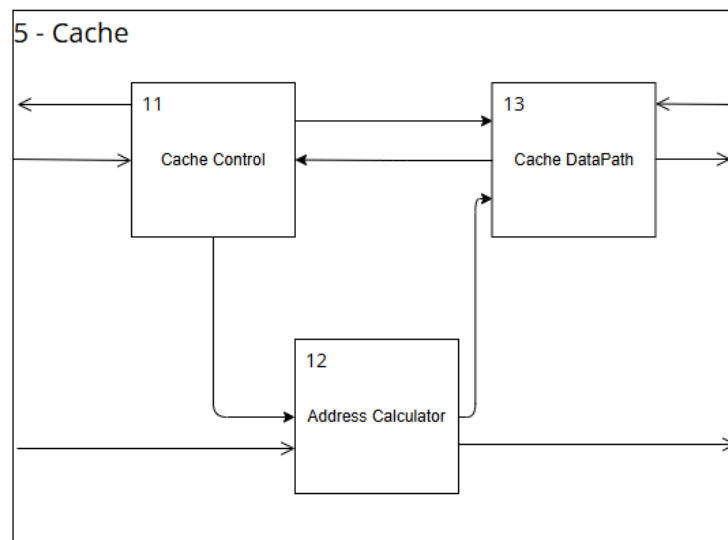


Figure 7: Cache structure diagram

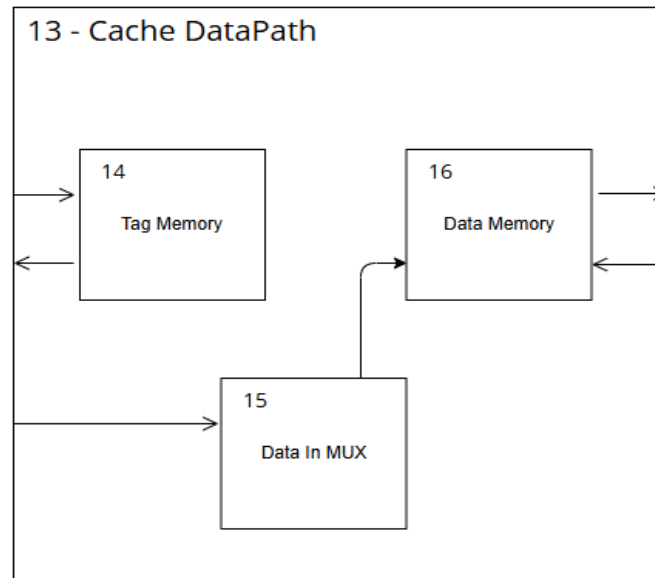


Figure 8: Cache DataPath architecture diagram

4.1 Hardware Description

The system was implemented in Verilog codes and schematic draws and loaded into a FPGA platform using the following hardware blocks:

- **DLX Based Core (x2):** Each core supports the DLX instruction set⁴.
- **Instruction Memory (x2):** One per core to eliminate fetch collisions.
- **Data Cache (x2):** Each core has a local data cache.
- **Main Memory:** Shared between cores, accessed via the bus.
- **Bus Controller:** Manages access scheduling and coherence signals.
- **MESI Controller:** FSM logic enforcing cache coherence for each core, as extension of the simple cache unit⁵.

In the following sub-sections, the different units, modules and components that we built and used during the project will be discussed and, when deemed necessary, an abstract figure will be attached, showing the modules' ports and connections⁶

⁴ Since the simple DLX processor itself is NOT part of the project rather a starting point and used as a "close box", its internal structure will not be addressed extensively in this paper. For more details regarding the DLX internal structure see G. Even, M. Markov & M. Medina (2015). *Computer Structure Lab Notes – Implementing a DLX processor on an FPGA*, Tel Aviv University. For more details about the DLX simplified architecture that it implements, see *Appendix B – Simplified DLX Instruction Set Architecture*

⁵ See footnote No. 3

⁶ More detailed elaboration about the different modules' ports, connections and I/O signals is presented in *Appendix A – Inputs-Outputs Tables*

Detailed explanation:

4.1.1 Main Memory: The main memory is implemented by the FPGA memory components in the physical implementation, and during simulations by the I/O Simul module, which was given in the ACSL course⁷.

Both consist of ordinary DRAM modules, and used as close boxes, the same they have been used during the ACSL course.

4.1.2 Bus Control: The bus control unit serves three main goals: data routing between the cores and the main memory, signal transmission and routing, and arbitration – managing the queue of the cores to memory access. To accomplish this the bus control receives all the data buses and signal lines from each of the above, holding in an internal memory the last core to address the main memory and perform the necessary logic computations to allow fast and accurate operation of the full system.

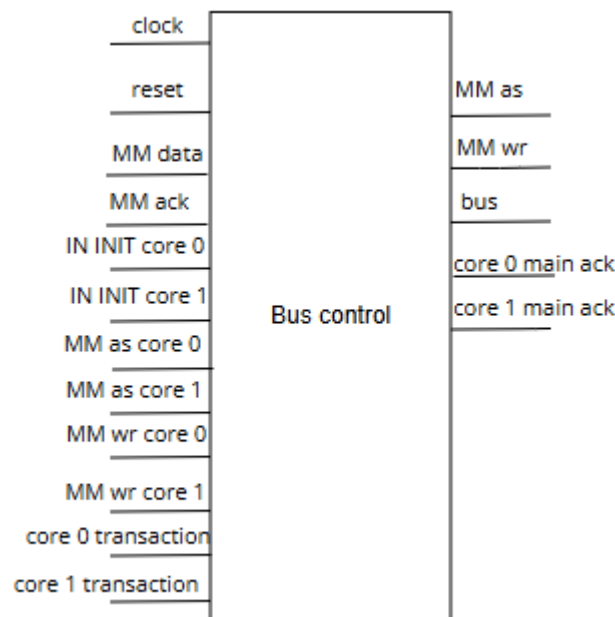


Figure 9: Bus Control Unit

⁷ Since both was not part of the work in the project and only used as close boxes, we felt there is no need to provide deeply detailed explanation and/or internal structure diagrams regarding them.

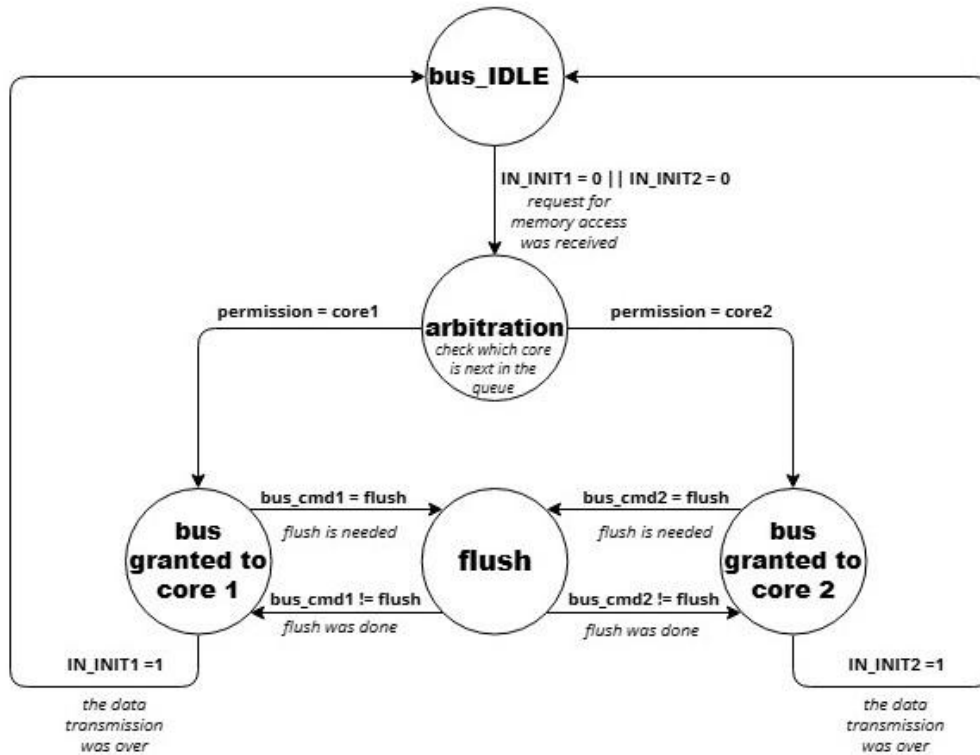


Figure 10: bus control state diagram

4.1.3 Core: This unit is a single core of the multicore processor architecture. The unit contain the DLX processor, the cache, the instruction memory, fetch state check unit, AS_N – DLX control signal DEMUX unit, DLX data in MUX unit, and the ACK signal MUX unit all detailed further on.

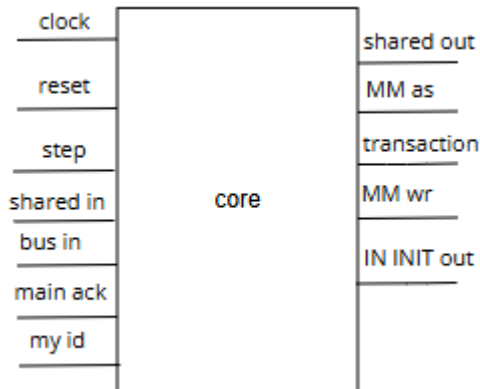


Figure 11: Core Unit

4.1.4 Fetch State Check: This unit checks the state of the DLX state machine and indicates if the DLX fetches an instruction.



Figure 12: Fetch State Check Unit

- 4.1.5 Cache: This unit is the cache memory of the core as explained in section 2. The unit contains the cache control, cache DATAPATH and address calculator unit detailed further on.

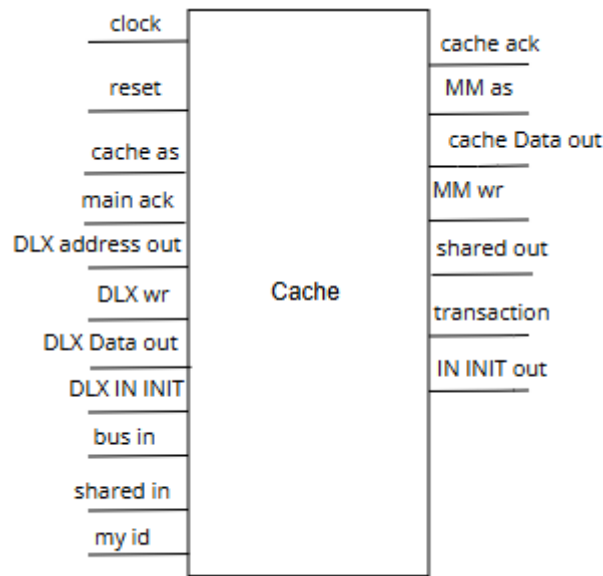


Figure 13: Cache Unit

- 4.1.6 AS_N – DLX control signal DEMUX: A DEMUX to route the AS_N signal of the DLX to the correct destination. If DLX fetches an instruction the signal goes to the instruction memory and if the DLX read/write data, the signal goes to the cache.




Figure 14: AS_N – DLX control signal DEMUX Unit

- 4.1.7 Instruction Memory: Pre-loaded with the compiled assembly code. Act as an instruction cache memory for the DLX.




Figure 15: Instruction Memory Unit

- 4.1.8 DLX data in MUX: A MUX for the DLX 'data in' input if the DLX fetch an instruction the data arrives from the instruction memory and if the DLX read data the data arrives from the cache.




Figure 16: DLX data in MUX Unit

- 4.1.9 ACK signal MUX: A MUX for the DLX ack signal if the DLX fetch an instruction the ack arrives from the instruction memory and if the DLX read/write data the ack arrives from the cache.




Figure 17: ACK signal MUX Unit

- 4.1.10 DLX: The DLX is a simple processor, implemented following a simplified version of the DLX architecture. It was implemented during the ACSL course⁸.

- 4.1.11 Cache Control: The cache control unit manages both data transmissions from and to the main memory, and access to the data saved inside the cache. This is done using a finite state machine, which is described in the diagram below (Fig. 19), and its decisions determined by the validity of the data inside the cache, the MESI state of the relevant block and additional information and requests which may will be received from the other cache.

It can perform either internal or external data access (depends if the block exists inside the cache and its MESI state, which determines its validity), data transmission to the main memory ("clearing block"), data transmission to the other cache ("flush", appears in purple in the diagram) or inform the other cache about a change in the status of certain block (if the block exists in both caches and it will be modified by 'write' instruction; appears in blue in the diagram). Since each block is 16-words size, each block transmission lasts 16 cycles, and the end of the cycle will be determined by internal counter (appears in italic form in the diagram).

⁸ See footnote No. 4

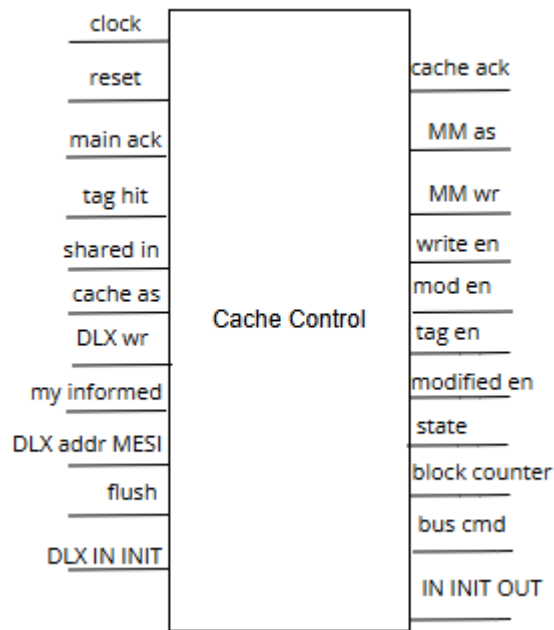


Figure 19: Cache Control Unit

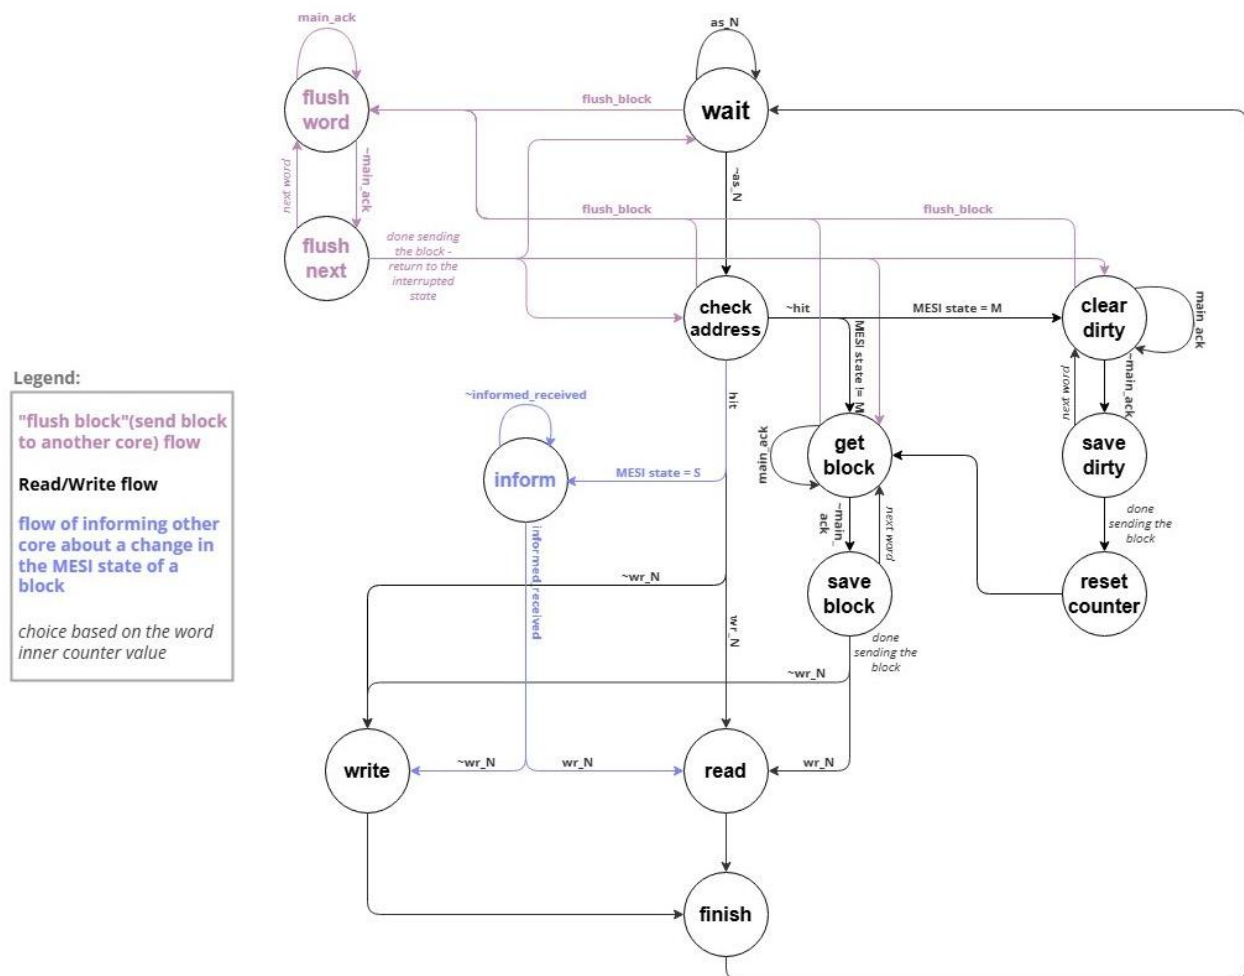


Figure 18: Cache control state diagram

4.1.12 Address Calculator: Address Calculator unit calculates the relevant address based on the required action. For example, when a full block is read or written the calculator outputs the 16 words addresses 1 by 1. Additional cases are flush, replacing modified block with new block and more. The calculator receives signals from the cache control and based on them perform the calculation.

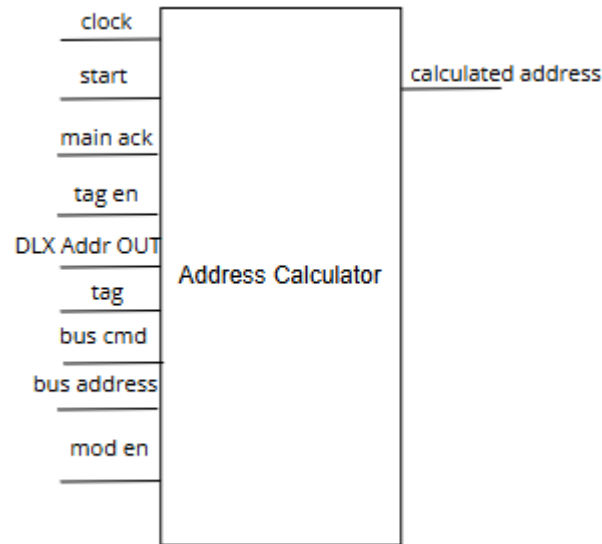


Figure 20: Address Calculator Unit

4.1.13 Cache DataPath: The Cache Datapath contains the units: Tag memory, Data in MUX and Data memory which are detailed further on.

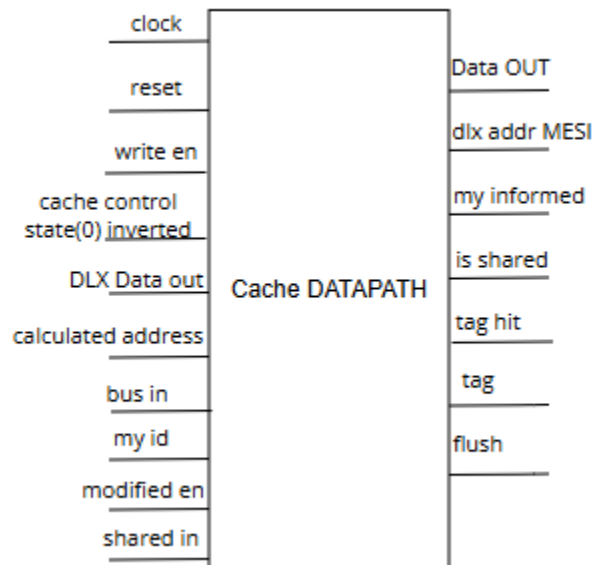


Figure 21: Cache DataPath Unit

4.1.14 Tag Memory: The tag memory unit contains the tags of the blocks that are in the cache and their MESI state. The unit "snoops" the bus and based on the transaction it detects the unit changes the state of the relevant block. The unit turns on the "shared" line when a block that is in the cache memory is requested by the other cache and if "flush" is needed the unit signals to the cache control.

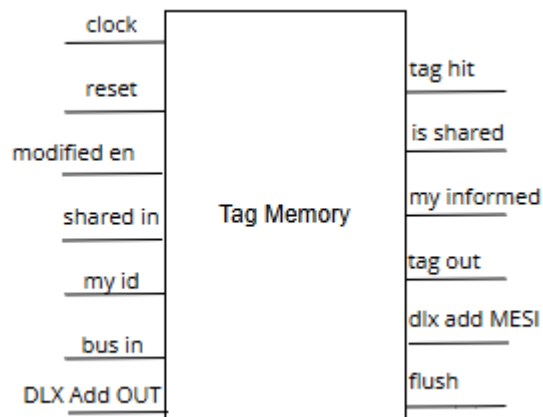


Figure 23: Tag Memory Unit

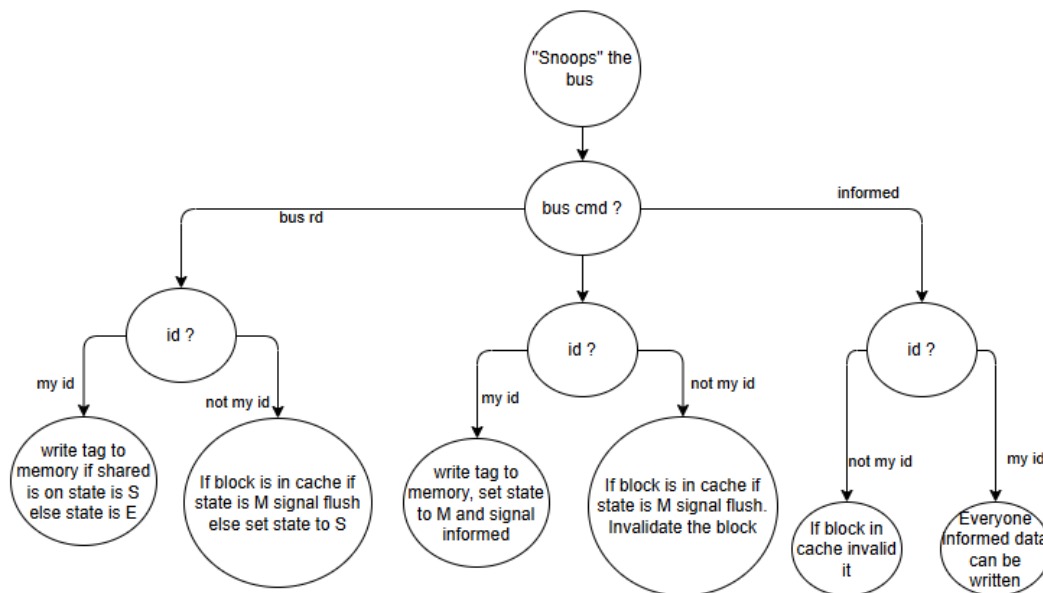


Figure 22: Tag Memory flow diagram

- 4.1.15 Data in MUX: Data in MUX is a 32bit MUX which based on the select signal chooses the data to be sent to the cache. When data from main memory or the other cache need to be written to the cache the select chooses 'bus data' input and in case the data is from the DLX the 'DLX data' input is chosen.




Figure 24: Data In MUX Unit

- 4.1.16 Data Memory: This unit is the memory of the cache. It contain up to 256 words divided into 16 blocks of data. Each clock cycle if 'write en' is on the word in 'Data in' input is written to the address based on the 'address' input and 'Data Out' output is the data that is in the requested address.




Figure 25: Cache Data Memory Unit

4.2 Software Description

Test programs were written in DLX assembly to exercise parallel access patterns, and demonstrate and validate coherence behavior under different parallelism conditions:

- **Bubble Sort** – Sorting a 32-words array from the main memory.
Each core sorting part of the array. This tests concurrent writing and enforces invalidations.
- **Vector Summing** – Addition of two 16-words vectors and save the result back in the main memory.
Independent sums are computed and shared, validating read-modify-write operations.
- **Basic image processing** – Importing an image, which is made of 16X16 array and performing a basic image processing function – segmentation by a given threshold value, e.g. 128.
Each core processes a block of image data in parallel, simulating independent access patterns.
- Additional code was written to validate the correctness of the design, by going through every transition possibly for the MESI protocol, every new function of the processor and performing additional tests.

All programs were compiled using RESA compilation software and loaded using Xilinx's ISE into the instruction memory units.

For each program we expected a significant speedup with minimal contention, particularly in memory-localized tasks

5. Analysis of Results

Functional Verification

All MESI transitions were verified through simulation and hardware traces. Cache states responded correctly to bus events. No inconsistencies in memory were observed.

Performance Metrics

After execution of all the programs on each of the processors (original simple DLX, DLX with cache extension and multicore DLX) the following results were recorded:

Table 2: Run time on each processor for different programs

Run time for program	Multicore[us]	Cache single-core[us]	Simple DLX[us]
Vector Addition	30.43782	46.43142	45.24023
Image Processing	172.7642	344.81202	427.14574
Bubble Sort	296.00655	695.87154	1120.2184

The results are also shown in the graph below for visual information and easy comparison between the different processors:




Figure 26: Programs runtime graph

By analysis of the results in the table above, the speedup in performance that was received is calculated below. As more complete analysis, the speedup was calculated not only between the original DLX and the multicore DLX, but between each of the processors above – for each program and the total average.

Table 3: Speedup ratio for any two processors

Speedup per program	DLX/multicore	DLX/cache single-core	cache single-core/multicore
Vector Addition	1.486316366	0.974345	1.525452
Image Processing	2.472420444	1.238779	1.995853
Bubble Sort	3.784437878	1.609806	2.350865
Average	2.581058229	1.27431	1.95739

A graphic representation and comparison can be seen in the graph below:




Figure 27: Speedup ratio graph

Several insights can be seen from the results above:

- Although some speedup was achieved by the addition of the cache to a single core processor, the majority of the improvement was due to the parallelism, which took advantage of the ability not only to do multiple computations in the same time but also to perform calculations by each core, when the other core might be IDLE and waiting for data transactions with the main memory to be completed
- The greatest improvement can be received in workloads with high amount of memory access actions, mostly in close addresses and local areas in the memory, with high number of computations and inter-core actions such as the bubble sort program.
- In programs with minimal number of computations, such as the vector addition, the speedup is less significant due to the cache structure, which requires the transmission of an entire block each time instead of a single word, and very little amount of calculation done internally in the cores.

Although several improvements contributed to superiority of the multicore processor in performance, it is clear that the main factor is parallelism, and its significance is well appreciated in the modern hardware industry and therefore its presence in every modern computer.

The execution of the different programs was not only achieved using the simulator, but on the physical implementation of the design, using the FPGA board. For example, the results of the bubble sort program can be seen below, as it was received using the RESA program.

The unsorted array - before the execution:

Memory Dump								
ADDRESS	VALUES							
0x00000000	0x000000a2	0x000000c4	0x000000e9	0x000000b3	0x00000071	0x00000087	0x000000f7	0x00000034
0x00000008	0x00000035	0x00000079	0x00000096	0x000000ea	0x000000d1	0x00000075	0x0000000b	0x000000e0
0x00000010	0x000000dd	0x0000005b	0x0000004f	0x000000ec	0x000000de	0x00000081	0x000000c9	0x000000af
0x00000018	0x000000d9	0x000000c3	0x00000062	0x0000009a	0x000000a5	0x0000001a	0x000000b8	0x00000044

Figure 28: Execution example - the array before sorting

The sorted array - after the execution:

Memory Dump								
ADDRESS	VALUES							
0x00000030	0x0000000b	0x0000001a	0x00000034	0x00000035	0x00000044	0x0000004f	0x0000005b	0x00000062
0x00000038	0x00000071	0x00000075	0x00000079	0x00000081	0x00000087	0x00000096	0x0000009a	0x000000a2
0x00000040	0x000000a5	0x000000af	0x000000b3	0x000000b8	0x000000c3	0x000000c4	0x000000c9	0x000000d1
0x00000048	0x000000d9	0x000000dd	0x000000de	0x000000e0	0x000000e9	0x000000ea	0x000000ec	0x000000f7

Figure 29: Execution example - the array after sorting

Comparison to other Cache Coherence Protocols: MESI vs MOESI, MSI, and Dragon

Modern multiprocessors rely on cache coherence protocols to ensure consistency between private caches. The most widely known is MESI, offering a good balance of efficiency and complexity. More advanced protocols like MOESI optimize memory traffic by allowing dirty sharing, while simpler ones like MSI trade performance for ease of implementation. Dragon, a write-update protocol, is suited for systems with frequent sharing and minimizes invalidations. Below is a compact comparison across key features.

Table 4: Different coherence protocol comparison

Feature	MESI	MOESI	MSI	Dragon
States	4: Modified, Exclusive, Shared, Invalid	5: MESI + Owned	3: Modified, Shared, Invalid	4: Modified, Shared, Exclusive, Shared-Modified
Write Policy	Write-back	Write-back	Write-back	Write-back + Write-through
Shared Read	S state	S/O states	S state	S/Sm states
Exclusive Ownership	E state	E/O states	None	E state
Dirty Sharing	None	O state	None	Sm state
Bus Traffic Efficiency	Medium	High	Low	High
Write Miss Cost	High	Moderate	High	Moderate
Read Miss Cost	Medium	Low (can get from O state)	High	Low
Implementation Complexity	Moderate	High	Low	High
Use Case	General purpose	High-performance CPUs	Simpler designs	Systems with frequent sharing

It can be seen that although more advanced protocols like MOESI and Dragon offer better performance, the MESI protocol strikes an effective balance between simplicity and performance

6. Conclusions and Further Work

Achievements

This project successfully extended a basic single-core DLX processor into a functional dual-core architecture with coherent memory access. The addition of local instruction and data caches to each core, combined with a carefully designed MESI-based coherence protocol, significantly improved overall system performance while maintaining data correctness. The design was thoroughly verified through simulation and deployed on FPGA hardware, demonstrating practical feasibility and robustness.

Improvements Observed

The evaluation across multiple benchmark programs—such as bubble sort, vector addition, and image processing—revealed substantial performance gains, particularly in memory-intensive tasks. The integration of parallel execution and local caching proved highly effective, with speedups of up to 3.78× over the baseline DLX processor. While caching alone provided modest improvements, the real impact stemmed from parallelism and efficient memory management (such as separated instructions memory unit, to reduce bottle neck in memory access).

Further Work

Several potential enhancements can further elevate the system's capabilities. These include both improvements in single-core level such as pipelining the DLX core for improved instruction throughput, introducing dynamic scheduling for out-of-order execution, and integrating branch prediction mechanisms, and on the multicore level – exploring advanced coherence protocols such as MOESI or Dragon, as well as dynamic cache replacement strategies and OS-level interrupt handling.

Implementing those improvements could enhance scalability and efficiency to provide significantly better performance of each core and the system as a whole.

Takeaways

This project illustrates the transformative value of multiprocessing in computer architecture, even in small-scale systems. By combining hardware-based parallelism with thoughtful cache and memory management, the system achieves substantial performance gains with manageable complexity. The results reflect trends in modern processor design and underscore the relevance of coherence protocols in maintaining data integrity across cores.

7. Project Documentation

The project was mainly documented using GitHub repository, using different branches for each part of the project (by stages).

Project files include:

- Verilog source files for cores, caches, and control units.
- Testbenches and simulation scripts, written in Verilog.
- FPGA .bit files. Used to load the design into the FPGA board.

Documentation and source code repository: Final-Project GitHub repository

User guide

Simulation:

- save two complied assembly code file name imem0 and imem1 for each core as .data file in the directory of all the project files named HOME_VER. Maximum of 128 instructions per file
- In the same directory save a sram.data file contains the initial main memory data in eight hexadecimal numbers, when empty lines considered as zeros. The file has a maximum of 1024 lines, when each line is a single word
- open HOME_VER.xise and run the testbench via xilinx platform and the waveform window will open.
- Review the waveforms as needed.

FPGA:

- save two complied assembly code file name imem0 and imem1 for each core as .data file in the directory of all the project files named SOURCE_VER. Maximum of 128 instructions per file.
- Create a .cod file via RESA software with the main memory initial data.
- Open SOURCE_VER.xise and create a bit file for the project.
- Using the RESA software upload the bit file and the .cod file to the FPGA
- Use continuous mode and let the programs run to halt.

8. References

1. Hennessy, J.L., & Patterson, D.A. (2017). *Computer Architecture: A Quantitative Approach*. Morgan Kaufmann.
2. Shen J. P., Lipasti M. K., "Modern Processor Design–Fundamentals of Superscalar Processors", McGraw-Hill, 1st ed. ,2005
3. Wikipedia contributors. (2025). *MESI protocol*. In Wikipedia, The Free Encyclopedia. https://en.wikipedia.org/wiki/MESI_protocol
4. G. Even, M. Markov & M. Medina (2015). *Computer Structure Lab Notes – Implementing a DLX processor on an FPGA*, Tel Aviv University.
5. *Multiprocessors Cache Coherence* and *Cache Coherence Examples* course handouts, from Computer Architecture course notes. G. Oxman (2024), Tel Aviv University

Appendix A. Inputs outputs tables

Bus control:

Table 5: Bus Control Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset
MM data	Input	32	Main memory data out
MM ack	Input	1	Main memory ack signal
IN INIT core 0	Input	1	IN INIT signal from core 0
IN INIT core 1	Input	1	IN INIT signal from core 1
MM as core 0	Input	1	AS_N signal received from core 0 dedicated for the main memory
MM as core 1	Input	1	AS_N signal received from core 1 dedicated for the main memory
MM wr core 0	Input	1	WR_N signal received from core 0 dedicated for the main memory
MM wr core 1	Input	1	WR_N signal received from core 1 dedicated for the main memory
Core 0 transaction	Input	69	Core 0 requested transaction
Core 1 transaction	Input	69	Core 1 requested transaction
MM as	Output	1	AS_N signal to main memory
MM wr	Output	1	WR_N signal to main memory
bus	Output	69	Bus data
Core 0 main ack	Output	1	ACK signal to core 0
Core 1 main ack	Output	1	ACK signal to core 1

DLX Core:

Table 6: Core Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset
step	Input	1	Step enable for DLX
Shared in	Input	1	The other core shared signal
Bus in	Input	69	Bus data
Main ack	Input	1	Ack from bus control
My id	Input	1	Core id
Shared out	Output	1	Shared signal to the other core
MM as	Output	1	Connection between cache 'MM as' output to bus control 'MM as core' input
transaction	Output	69	Requested transaction

MM wr	Output	1	Connection between cache 'MM wr' output to bus control 'MM wr core' input
IN INIT out	Output	1	Connection between cache 'IN INIT out' output to bus control 'IN INIT core' input

Fetch state check:

Table 7: Fetch State Check Unit Input Output table

Name	Input/ Output	Size [bits]	Description
DLX CTRL state	Input	5	DLX control state machine state.
fetch	Output	1	Fetch = 1 if DLX control state is fetch

Cache:

Table 8: Cache Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset
Cache as	Input	1	Connection between AS_N DEMUX 'cache as' output to the cache control 'cache as' input.
Main ack	Input	1	Connection between bus control 'Core main ack' output to cache control 'main ack' input.
DLX address out	Input	32	Connection between DLX 'address out' output to address calculator 'DLX Add OUT' input.
DLX wr	Input	1	Connection between DLX 'wr_n' output to cache control 'DLX wr' input
DLX data out	Input	32	Connection between DLX 'data out' output to cache DATAPATH 'DLX data out' input
DLX IN INIT	Input	1	Connection between DLX 'IN INIT' output to cache control 'DLX IN INIT' input
Bus in	Input	69	Bus data
Shared in	Input	1	The other core shared signal
My id	Input	1	Core id
Cache ack	Output	1	Connection between cache control 'cache ack' output to ACK signal MUX 'cache ack' input.
MM as	Output	1	Connection between cache control 'MM as' output to bus control 'MM as core' input.
Cache data out	Output	32	Connection between cache DATAPATH 'data out' output to DLX data in MUX 'cache data out' input
MM wr	Output	1	Connection between cache control 'MM wr' output to bus control 'MM wr core' input.
Shared out	Output	1	Connection between cache DATAPATH 'is shared' output to the other core 'shared in' input
transaction	Output	69	Transaction requested.
IN INIT out	Output	1	Connection between cache control 'IN INIT out' output to bus control 'IN INIT core' input

AS_N – DLX control signal DEMUX:

Table 9: AS_N – DLX control signal DEMUX Unit Input Output table

Name	Input/ Output	Size [bits]	Description
fetch	Input	1	Selects where the as signal needs to go.
DLX as	Input	1	DLX as signal. Signaling of an incoming address to be read/write from/to
Cache as	Output	1	Cache as = DLX as if fetch = 0 else equals 1
Imem as	Output	1	Imem as = DLX as if fetch = 0 else equals 1

Instruction Memory:

Table 10: Instruction Memory Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
My id	Input	1	Core id
Imem as	Input	1	Signals to imem that an instruction is wished to be read
DLX address out	Input	7	The pc of the instruction that wished to be read
Imem data out	Output	32	The instruction that the DLX wishes to read
Imem ack	Output	1	Ack signal to signal that the instruction is ready to be read

DLX data in MUX:

Table 11: DLX data in MUX Unit Input Output table

Name	Input /Output	Size [bits]	Description
fetch	Input	1	Fetch = 0 => DLX data in = cache data out Fetch = 1 => DLX data in = imem data out
Cache data out	Input	32	Cache data out
Imem data out	Input	32	Imem data out
DLX Data in	Output	32	The data the DLX wish to read base of if it is an instruction or data from the cache.

ACK signal MUX:

Table 12: ACK signal MUX Unit Input Output table

Name	Input/ Output	Size [bits]	Description
fetch	Input	1	Fetch = 0 => DLX ack = cache ack Fetch = 1 => DLX ack = imem ack
Cache ack	Input	1	Ack signal from cache control indicates to DLX that the data is ready to be read/written
Imem ack	Input	1	Ack signal from imem indicates to DLX that the instruction is ready to be read
DLX ack	Output	1	Ack signal to DLX based on if it reads instruction or reads/write data.

Cache control:

Table 13: Cache Control Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset
Main ack	Input	1	Ack signal from the bus control
Tag hit	Input	1	Signal from tag memory if the block is in cache or not.
Shared in	Input	1	The other core shared signal
Cache as	Input	1	As signal from DLX indicates an address is set to read/write data.
DLX wr	Input	1	Signal if the DLX write or read
My informed	Input	1	Signal from tag memory that informed transaction of my core is detected on the bus.
DLX addr MESI	Input	2	The MESI state of the block of address requested by the DLX
flush	Input	1	Signal from tag memory that a flush of a block is needed
DLX IN INIT	Input	1	DLX IN INIT signal
Cache ack	Output	1	Cache control signal to DLX that the data is ready to be read or the data was written
MM as	Output	1	Cache control signals to bus control of an address arriving for read/write
MM wr	Output	1	Cache control signals to bus control if read or write
Write en	Output	1	Cache control signals to data memory to write the data.
Mod en	Output	1	Indicates address calculator if a block is read/written or a single word
Tag en	Output	1	Indicates address calculator where to take the tag from (current in cache or requested)
Modified en	Output	1	Cache control signals to tag memory to set the block MESI state to M

state	Output	4	The state of the cache control state machine
Block counter	Output	5	When block counter equals 0 the address calculator starts calculating the address when a full block is read/written
Bus cmd	Output	3	One of the indicators for the address calculator.
IN INIT out	Output	1	IN INIT output to bus control

Address calculator:

Table 14: Address Calculator Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
start	Input	1	NOR of cache control 'block counter' output bits. Signals to start calculate the addresses when a whole block is read/written.
Main ack	Input	1	Ack from main memory signals that the data of the current address was read/written.
Tag en	Input	1	Indicates address calculator where to take the tag from (current in cache or requested)
DLX Add OUT	Input	32	DLX address out
tag	Input	24	Tag of the block of the address on the bus given from tag memory.
Bus cmd	Input	3	Bus command requested by the cache control determines the calculation of the address
Bus address	Input	32	The address that is on the bus
Mod en	Input	1	Indicates address calculator if a block is read/written or a single word
Calculated address	Output	32	Calculated address

Cache DATAPATH:

Table 15: Cache DataPath Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset
Write en	Input	1	Connection between cache control 'write en' output to data memory 'write en' input
Cache control state(0) inverted	Input	1	Connection between Bit 0 of cache control state inverted to Data in MUX 'select' input
DLX Data out	Input	32	Connection between DLX 'data out' output to Data in MUX 'DLX data' input
Calculated address	Input	32	Connection between address calculator 'Calculated address' output to Tag memory 'DLX Add OUT' input and Data memory 'Address' input
Bus in	Input	69	Bus data

My id	Input	1	Core id
Modified en	Input	1	Connection between cache control 'modified en' output to Tag memory 'modified en' input
Shared in	Input	1	The other core shared signal
Data out	Output	32	Cache data out
DLX addr MESI	Output	2	Connection between Tag memory 'DLX addr MESI' output to cache control 'DLX addr MESI' input
My informed	Output	1	Connection between Tag memory 'my informed' output to cache control 'my informed' input
Is shared	Output	1	Connection between Tag memory 'is shared' output to the other core 'shared in' input
Tag hit	Output	1	Connection between Tag memory 'tag hit' output to cache control 'tag hit' input
tag	Output	24	Connection between Tag memory 'tag out' output to address calculator 'tag' input
flush	Output	1	Connection between Tag memory 'flush' output to cache control 'flush' input

Tag Memory:

Table 16: Tag Memory Unit Input Output table

Name	Input/ Output	Size [bits]	Description
clock	Input	1	Clock signal
reset	Input	1	Global reset all blocks set to invalid
Modified en	Input	1	Cache control signals to set the block to state modified
Shared in	Input	1	The other core shared signal
My id	Input	1	Core id
Bus in	Input	37	Bus cmd, bus origid and bus address that are read from the bus
DLX Add OUT	Input	32	The Address the DLX wishes to write/read to/from after address calculator calculation
Tag hit	Output	1	Signals cache control the block is in the cache
Is shared	Output	1	On if the block of the address that read from the bus is in the cache
My informed	Output	1	Signals cache control that an informed transaction with my id is read on the bus
Tag out	Output	24	The current tag of the block of the address the read from the bus
DLX add MESI	Output	2	The MESI state of the block of address requested by the DLX
flush	Output	1	Signals cache control that a flush is required

Data in MUX:

Table 17: Data in MUX Unit Input Output table

Name	Input/ Output	Size [bits]	Description
select	Input	1	Select = 0 => Data in = bus data Select = 1 => Data in = DLX data
Bus data	Input	32	Data that is written on the bus
DLX data	Input	32	Output data of the DLX
Data In	Output	32	Cache's data memory data in

Data Memory:

Table 18: Cache Data Memory Unit

Name	Input/ Output	Size [bits]	Description
Clock	Input	1	Clock signal
Write_EN	Input	1	Write 'Data In' to memory in 'Address' when enabled
Address	Input	8	Address in data memory array the DLX request to read/write from/to. after address calculator calculation
Data In	Input	32	Data to be written
Data Out	Output	32	Data in memory at 'Address'

Appendix B. The Simplified DLX Instruction Set Architecture⁹

We list below the instruction set of the simplified DLX. Note that *imm* denotes the value of the immediate field in an I-Type-Instruction. *sext(imm)* denotes the 2's complement sign extension of *imm* to 32 bits, R(N) denotes value of the register number N, M(A) denotes value of the memory address A.

- Load/Store-instructions:

Load/Store				Semantics
Lw	RD	RS1	imm	$R(RD) := M(\text{sext}(\text{imm}) + R(RS1))$
Sw		RD RS1	imm	$M(\text{sext}(\text{imm}) + R(RS1)) := R(RD)$

- Immediate-instructions:

Instruction				Semantics
Addi	RD	RS1	imm	$R(RD) := R(RS1) + \text{sext}(\text{imm})$

- Shift-/Compute-Instructions:

Instruction				Semantics
slli	RD	RS1		$R(RD) := R(RS1) \ll 1$
srl	RD	RS1		$R(RD) := R(RS1) \gg 1$
add	RD	RS1	RS2	$R(RD) := R(RS1) + R(RS2)$
sub	RD	RS1	RS2	$R(RD) := R(RS1) - R(RS2)$
and	RD	RS1	RS2	$R(RD) := R(RS1) \wedge R(RS2)$
or	RD	RS1	RS2	$R(RD) := R(RS1) \vee R(RS2)$
xor	RD	RS1	RS2	$R(RD) := R(RS1) \oplus R(RS2)$

- Test-Instruction

Instruction				Semantics
Sreli	RD	RS1	imm	$R(RD) := 1$, if condition is satisfied $R(RD) := 0$ otherwise
If <i>rel</i> =lt				test if $R(RS1) < \text{sext}(\text{imm})$
If <i>rel</i> =eq				test if $R(RS1) = \text{sext}(\text{imm})$
If <i>rel</i> =gt				test if $R(RS1) > \text{sext}(\text{imm})$
If <i>ref</i> =le .				test if $R(RS1) \leq \text{sext}(\text{imm})$
If <i>ref</i> =ge .				test if $R(RS1) \geq \text{sext}(\text{imm})$
If <i>ref</i> =ne .				test if $R(RS1) \neq \text{sext}(\text{imm})$

- Jump-instructions

Instruction				Semantics
beqz		RS1	imm	$PC := PC + 1 + \text{sext}(\text{imm})$, if $R(RS1) = 0$
bnez		RS1	imm	$PC := PC + 1$, if $R(RS1) \neq 0$
				$PC := PC + 1 + \text{sext}(\text{imm})$, if $R(RS1) \neq 0$
jr		RS1		$PC := R(RS1)$
Jalr		RS1		$R(31) := PC + 1$; $PC := R(RS1)$

⁹ From G. Even, M. Markov & M. Medina (2015). *Computer Structure Lab Notes – Implementing a DLX processor on an FPGA*, Tel Aviv University, pages 39-40.

- Miscellaneous-instructions

Instruction	Semantics
special-nop	causes transition the DLX to Init/Fetch states
halt	causes transition the DLX to HALT state